

Latent AI

Enterprise language models that outperform
GPT-4 at 1/10th the cost.

Series A · Confidential · 2025

THE PROBLEM

Frontier models are too expensive,
too slow, and too risky for enterprise.

\$0.06

per 1K tokens

GPT-4 output pricing makes high-volume enterprise use cases like document processing and customer service prohibitively expensive at scale.

800ms

median latency

Frontier model response times are too slow for real-time applications. Financial services, healthcare, and legal workflows demand sub-200ms.

Zero

data control

Sending proprietary enterprise data to third-party APIs creates compliance risk. Regulated industries cannot use shared model infrastructure.

BENCHMARKS

Latent-70B: frontier performance at a fraction of the cost.

Latent-70B is a 70-billion parameter model trained on 4.2 trillion tokens of curated enterprise data. It runs on a single 8xH100 node and serves at \$0.006 per 1K output tokens.

MMLU Score

Latent-70B achieves 86.4% on MMLU 5-shot, compared to GPT-4's published score of 86.4% and Llama 3 70B's 79.5%. Matching frontier performance at 10x lower serving cost.

Enterprise Tasks

Scores 91.2% on LegalBench, outperforming GPT-4 (84.7%) and Claude 3 Opus (87.1%) on legal reasoning tasks — the first open-weight model to surpass closed-source models on this benchmark.

Inference Speed

Generates 142 tokens/second on 8xH100 using our custom serving stack with speculative decoding, compared to 38 tokens/second for standard vLLM serving of Llama 3 70B.

Context Length

Supports 256K token context window with less than 5% degradation in needle-in-a-haystack retrieval accuracy at 200K+ tokens, validated on our internal long-context benchmark.

Our world-class research team includes former scientists from Google DeepMind, Meta FAIR, and OpenAI.

TRACTION

Enterprise adoption is accelerating.

28

Enterprise Customers

Including 4 Fortune 500 companies deploying in production

\$6.8M

ARR

Growing 32% month-over-month since launching 8 months ago

2.1B

Tokens Served Daily

Across legal, financial services, and healthcare verticals

99.97%

Uptime SLA

Self-hosted deployments with zero data leaving customer VPC

Latent AI

Raising \$32M Series A to build the enterprise
language model infrastructure layer.

founders@latent.ai

latent.ai